

```
In [1]: cd /home/nv/Downloads
```

```
/home/nv/Downloads
```

```
In [2]: import keras  
print keras.__version__
```

```
1.1.1
```

```
In [3]: import theano
```

```
In [5]: # Load and Plot the IMDB dataset
import numpy
from keras.datasets import imdb
from matplotlib import pyplot
# load the dataset
(X_train, y_train), (X_test, y_test) = imdb.load_data()
X = numpy.concatenate((X_train, X_test), axis=0)
y = numpy.concatenate((y_train, y_test), axis=0)
# summarize size
print("Training data: ")
print(X.shape)
print(y.shape)
# Summarize number of classes
print("Classes: ")
print(numpy.unique(y))
# Summarize number of words
print("Number of words: ")
print(len(numpy.unique(numpy.hstack(X))))
# Summarize review length
print("Review length: ")
result = map(len, X)
print("Mean %.2f words (%f)" % (numpy.mean(result), numpy.std(result)))
# plot review length as a boxplot and histogram
pyplot.subplot(121)
pyplot.boxplot(result)
pyplot.subplot(122)
pyplot.hist(result)
pyplot.show()
```

Training data:

(50000,)

(50000,)

Classes:

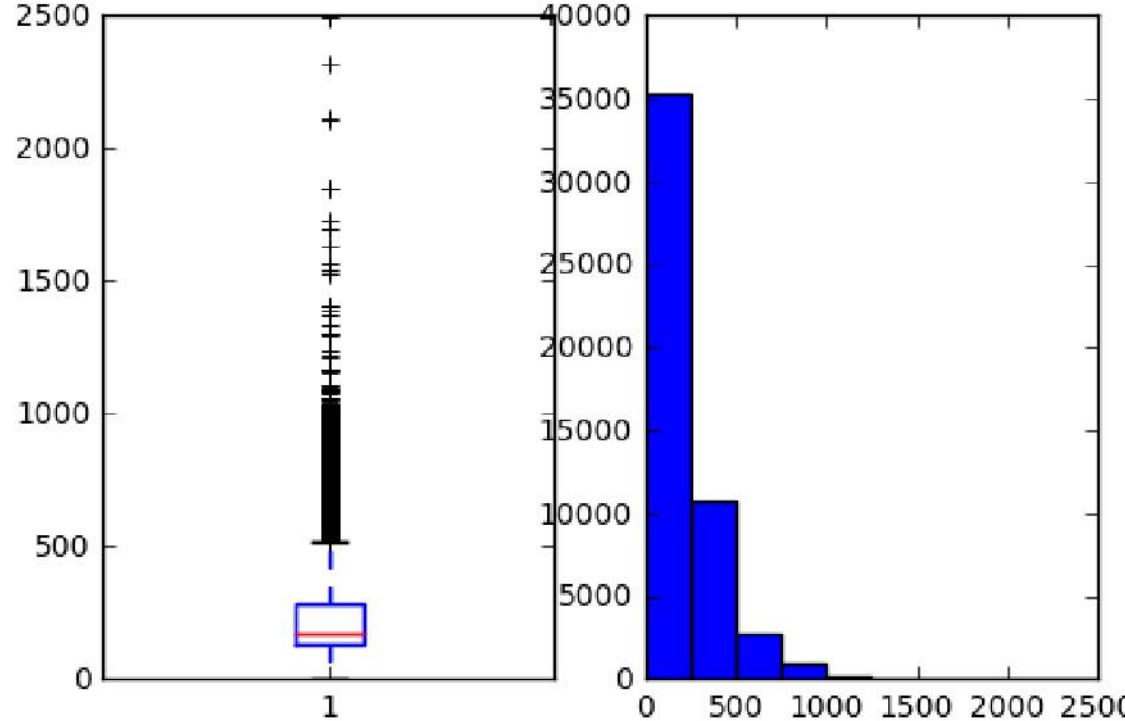
[0 1]

Number of words:

88585

Review length:

Mean 234.76 words (172.911495)



```
In [6]: # MLP for the IMDB problem, optimizer='adam', nb_epoch =10, input_length(max_words)=500
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
import matplotlib.pyplot as plt
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
max_words = 500
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=10, batch_size=128, verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_1 (Embedding)	(None, 500, 32)	160000	embedding_input_1[0][0]
flatten_1 (Flatten)	(None, 16000)	0	embedding_1[0][0]
dense_1 (Dense)	(None, 250)	4000250	flatten_1[0][0]

dense_2 (Dense)	(None, 1)	251	dense_1[0][0]
=====			
Total params: 4160501			
<hr/>			
None			
Train on 25000 samples, validate on 25000 samples			
Epoch 1/10			
25000/25000 [=====] - 55s - loss: 0.5366 - acc: 0.6864 - val_loss: 0.2966 - val_acc: 0.8758			
Epoch 2/10			
25000/25000 [=====] - 56s - loss: 0.2048 - acc: 0.9208 - val_loss: 0.2982 - val_acc: 0.8764			
Epoch 3/10			
25000/25000 [=====] - 58s - loss: 0.0725 - acc: 0.9794 - val_loss: 0.3935 - val_acc: 0.8655			
Epoch 4/10			
25000/25000 [=====] - 66s - loss: 0.0153 - acc: 0.9968 - val_loss: 0.4752 - val_acc: 0.8633			
Epoch 5/10			
25000/25000 [=====] - 55s - loss: 0.0025 - acc: 0.9999 - val_loss: 0.5428 - val_acc: 0.8669			
Epoch 6/10			
25000/25000 [=====] - 57s - loss: 7.3811e-04 - acc: 1.0000 - val_loss: 0.5797 - val_acc: 0.8674			
Epoch 7/10			
25000/25000 [=====] - 64s - loss: 4.1734e-04 - acc: 1.0000 - val_loss: 0.6054 - val_acc: 0.8681			
Epoch 8/10			
25000/25000 [=====] - 65s - loss: 2.7427e-04 - acc: 1.0000 - val_loss: 0.6274 - val_acc: 0.8681			
Epoch 9/10			
25000/25000 [=====] - 65s - loss: 1.9458e-04 - acc: 1.0000 - val_loss: 0.6472 - val_acc: 0.8684			
Epoch 10/10			
25000/25000 [=====] - 65s - loss: 1.4375e-04 - acc: 1.0000 - val_loss: 0.6633 - val_acc: 0.8684			
Accuracy: 86.84%			

As we can see this model overfits very quickly so i will change it very few training epochs and since there is a lot of data so i will keep batch size same as 128

```
In [7]: # MLP for the IMDB problem,optimizer='adam',input_length (max_words)=500, nb_epoch=5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
import matplotlib.pyplot as plt
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
max_words = 500
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128, verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_2 (Embedding)	(None, 500, 32)	160000	embedding_input_2[0][0]
flatten_2 (Flatten)	(None, 16000)	0	embedding_2[0][0]
dense_3 (Dense)	(None, 250)	4000250	flatten_2[0][0]
dense_4 (Dense)	(None, 1)	251	dense_3[0][0]
Total params: 4160501			

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 58s - loss: 0.5366 - acc: 0.6864 - val\_loss: 0.2966 - val\_acc: 0.8758

Epoch 2/5

25000/25000 [=====] - 58s - loss: 0.2048 - acc: 0.9208 - val\_loss: 0.2982 - val\_acc: 0.8764

Epoch 3/5

25000/25000 [=====] - 55s - loss: 0.0725 - acc: 0.9794 - val\_loss: 0.3935 - val\_acc: 0.8655

Epoch 4/5

25000/25000 [=====] - 48s - loss: 0.0153 - acc: 0.9968 - val\_loss: 0.4752 - val\_acc: 0.8633

Epoch 5/5

25000/25000 [=====] - 61s - loss: 0.0025 - acc: 0.9999 - val\_loss: 0.5428 - val\_acc: 0.8669

Accuracy: 86.69%

```
In [8]: # MLP for the IMDB problem for input length 1000 (max_words),optimizer='adam',nb_epoch = 5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
import matplotlib.pyplot as plt
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
max_words = 1000
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_3 (Embedding)	(None, 1000, 32)	160000	embedding_input_3[0] [0]
flatten_3 (Flatten)	(None, 32000)	0	embedding_3[0] [0]
dense_5 (Dense)	(None, 250)	8000250	flatten_3[0] [0]
dense_6 (Dense)	(None, 1)	251	dense_5[0] [0]
Total params: 8160501			

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 109s - loss: 0.4993 - acc: 0.7350 - val\_loss: 0.3011  
- val\_acc: 0.8716

Epoch 2/5

25000/25000 [=====] - 105s - loss: 0.1888 - acc: 0.9294 - val\_loss: 0.2982  
- val\_acc: 0.8744

Epoch 3/5

25000/25000 [=====] - 103s - loss: 0.0619 - acc: 0.9826 - val\_loss: 0.3901  
- val\_acc: 0.8668

Epoch 4/5

25000/25000 [=====] - 130s - loss: 0.0099 - acc: 0.9985 - val\_loss: 0.4957  
- val\_acc: 0.8660

Epoch 5/5

25000/25000 [=====] - 109s - loss: 0.0019 - acc: 0.9999 - val\_loss: 0.5418  
- val\_acc: 0.8678

Accuracy: 86.78%

```
In [14]: # MLP for the IMDB problem for input length 500 (max_words),optimizer='rmsprop',nb_epoch = 5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
import matplotlib.pyplot as plt
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
max_words = 500
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_9 (Embedding)	(None, 500, 32)	160000	embedding_input_9[0] [0]
flatten_9 (Flatten)	(None, 16000)	0	embedding_9[0] [0]
dense_17 (Dense)	(None, 250)	4000250	flatten_9[0] [0]
dense_18 (Dense)	(None, 1)	251	dense_17[0] [0]
Total params: 4160501			

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 60s - loss: 0.4978 - acc: 0.7324 - val\_loss: 0.3045 - val\_acc: 0.8711

Epoch 2/5

25000/25000 [=====] - 45s - loss: 0.1985 - acc: 0.9220 - val\_loss: 0.3105 - val\_acc: 0.8723

Epoch 3/5

25000/25000 [=====] - 45s - loss: 0.0791 - acc: 0.9730 - val\_loss: 0.4628 - val\_acc: 0.8523

Epoch 4/5

25000/25000 [=====] - 44s - loss: 0.0233 - acc: 0.9928 - val\_loss: 0.5901 - val\_acc: 0.8502

Epoch 5/5

25000/25000 [=====] - 44s - loss: 0.0063 - acc: 0.9983 - val\_loss: 0.7891 - val\_acc: 0.8426

Accuracy: 84.26%

```
In [15]: # MLP for the IMDB problem for input length 1000 (max_words),optimizer='rmsprop',nb_epoch = 5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
import matplotlib.pyplot as plt
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
max_words = 1000
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_10 (Embedding)	(None, 1000, 32)	160000	embedding_input_10[0][0]
flatten_10 (Flatten)	(None, 32000)	0	embedding_10[0][0]
dense_19 (Dense)	(None, 250)	8000250	flatten_10[0][0]
dense_20 (Dense)	(None, 1)	251	dense_19[0][0]
Total params: 8160501			

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 102s - loss: 0.5048 - acc: 0.7457 - val\_loss: 0.3070  
- val\_acc: 0.8697

Epoch 2/5

25000/25000 [=====] - 102s - loss: 0.2030 - acc: 0.9189 - val\_loss: 0.3002  
- val\_acc: 0.8754

Epoch 3/5

25000/25000 [=====] - 97s - loss: 0.0939 - acc: 0.9656 - val\_loss: 0.3962 -  
val\_acc: 0.8658

Epoch 4/5

25000/25000 [=====] - 103s - loss: 0.0325 - acc: 0.9892 - val\_loss: 0.5644  
- val\_acc: 0.8560

Epoch 5/5

25000/25000 [=====] - 98s - loss: 0.0087 - acc: 0.9971 - val\_loss: 0.6809 -  
val\_acc: 0.8546

Accuracy: 85.46%

```
In [9]: # CNN for the IMDB problem with input size =max_words 500,optimizer='adam', nb_epoch=5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.convolutional import Convolution1D
from keras.layers.convolutional import MaxPooling1D
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
# pad dataset to a maximum review length in words
max_words = 500
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Convolution1D(nb_filter=32, filter_length=3, border_mode='same',
activation='relu'))
model.add(MaxPooling1D(pool_length=2))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,
verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

```
/home/nv/anaconda2/lib/python2.7/site-packages/keras/backend/theano_backend.py:1500: UserWarning: DEPRECATION: the 'ds' parameter is not going to exist anymore as it is going to be replaced by the parameter 'ws'.
    mode='max')
/home/nv/anaconda2/lib/python2.7/site-packages/keras/backend/theano_backend.py:1500: UserWarning: DEPRECATION: the 'st' parameter is not going to exist anymore as it is going to be replaced by the parameter 'stride'.
    mode='max')
/home/nv/anaconda2/lib/python2.7/site-packages/keras/backend/theano_backend.py:1500: UserWarning: DEPRECATION: the 'padding' parameter is not going to exist anymore as it is going to be replaced by the parameter 'pad'.
    mode='max')
```

Layer (type)	Output Shape	Param #	Connected to
embedding_4 (Embedding)	(None, 500, 32)	160000	embedding_input_4[0][0]
convolution1d_1 (Convolution1D)	(None, 500, 32)	3104	embedding_4[0][0]
maxpooling1d_1 (MaxPooling1D)	(None, 250, 32)	0	convolution1d_1[0][0]
flatten_4 (Flatten)	(None, 8000)	0	maxpooling1d_1[0][0]
dense_7 (Dense)	(None, 250)	2000250	flatten_4[0][0]
dense_8 (Dense)	(None, 1)	251	dense_7[0][0]

Total params: 2163605

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 74s - loss: 0.4857 - acc: 0.7164 - val\_loss: 0.2764 - val\_acc: 0.8864

Epoch 2/5

25000/25000 [=====] - 74s - loss: 0.2249 - acc: 0.9123 - val\_loss: 0.2675 - val\_acc: 0.8890

Epoch 3/5

25000/25000 [=====] - 74s - loss: 0.1680 - acc: 0.9359 - val\_loss: 0.2855 - val\_acc: 0.8841

Epoch 4/5

25000/25000 [=====] - 120s - loss: 0.1242 - acc: 0.9554 - val\_loss: 0.3293 - val\_acc: 0.8810

Epoch 5/5

25000/25000 [=====] - 125s - loss: 0.0825 - acc: 0.9729 - val\_loss: 0.3956 - val\_acc: 0.8749

Accuracy: 87.49%

```
In [13]: # CNN for the IMDB problem with input size =max_words 1000,optimizer='adam', nb_epoch = 5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.convolutional import Convolution1D
from keras.layers.convolutional import MaxPooling1D
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
# pad dataset to a maximum review length in words
max_words = 1000
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Convolution1D(nb_filter=32, filter_length=3, border_mode='same',
activation='relu'))
model.add(MaxPooling1D(pool_length=2))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,
verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_8 (Embedding)	(None, 1000, 32)	160000	embedding_input_8[0] [0]
convolution1d_5 (Convolution1D)	(None, 1000, 32)	3104	embedding_8[0] [0]
maxpooling1d_5 (MaxPooling1D)	(None, 500, 32)	0	convolution1d_5[0] [0]
flatten_8 (Flatten)	(None, 16000)	0	maxpooling1d_5[0] [0]
dense_15 (Dense)	(None, 250)	4000250	flatten_8[0] [0]
dense_16 (Dense)	(None, 1)	251	dense_15[0] [0]
<hr/>			
Total params: 4163605			

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 156s - loss: 0.4529 - acc: 0.7555 - val\_loss: 0.3227  
- val\_acc: 0.8612

Epoch 2/5

25000/25000 [=====] - 186s - loss: 0.2360 - acc: 0.9060 - val\_loss: 0.2768  
- val\_acc: 0.8850

Epoch 3/5

25000/25000 [=====] - 170s - loss: 0.1769 - acc: 0.9329 - val\_loss: 0.3257  
- val\_acc: 0.8713

Epoch 4/5

25000/25000 [=====] - 212s - loss: 0.1306 - acc: 0.9521 - val\_loss: 0.3263  
- val\_acc: 0.8785

Epoch 5/5

25000/25000 [=====] - 209s - loss: 0.0845 - acc: 0.9724 - val\_loss: 0.3977  
- val\_acc: 0.8707

Accuracy: 87.07%

```
In [11]: # CNN for the IMDB problem with input size =max_words 500 , optimizer='rmsprop',nb_epoch =5
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.convolutional import Convolution1D
from keras.layers.convolutional import MaxPooling1D
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
# pad dataset to a maximum review length in words
max_words = 500
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Convolution1D(nb_filter=32, filter_length=3, border_mode='same',
activation='relu'))
model.add(MaxPooling1D(pool_length=2))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,
verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_6 (Embedding)	(None, 500, 32)	160000	embedding_input_6[0][0]
convolution1d_3 (Convolution1D)	(None, 500, 32)	3104	embedding_6[0][0]
maxpooling1d_3 (MaxPooling1D)	(None, 250, 32)	0	convolution1d_3[0][0]
flatten_6 (Flatten)	(None, 8000)	0	maxpooling1d_3[0][0]
dense_11 (Dense)	(None, 250)	2000250	flatten_6[0][0]
dense_12 (Dense)	(None, 1)	251	dense_11[0][0]

Total params: 2163605

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 100s - loss: 0.4547 - acc: 0.7606 - val\_loss: 0.3823  
- val\_acc: 0.8306

Epoch 2/5

25000/25000 [=====] - 98s - loss: 0.2494 - acc: 0.8985 - val\_loss: 0.2613 -  
val\_acc: 0.8920

Epoch 3/5

25000/25000 [=====] - 98s - loss: 0.2039 - acc: 0.9195 - val\_loss: 0.2836 -  
val\_acc: 0.8806

Epoch 4/5

25000/25000 [=====] - 158s - loss: 0.1755 - acc: 0.9320 - val\_loss: 0.2996  
- val\_acc: 0.8814

Epoch 5/5

25000/25000 [=====] - 109s - loss: 0.1482 - acc: 0.9443 - val\_loss: 0.3039  
- val\_acc: 0.8852

Accuracy: 88.52%

```
In [12]: # CNN for the IMDB problem with input size =max_words 1000,optimizer='rmsprop'
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Flatten
from keras.layers.convolutional import Convolution1D
from keras.layers.convolutional import MaxPooling1D
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
# fix random seed for reproducibility
seed = 7
numpy.random.seed(seed)
# load the dataset but only keep the top n words, zero the rest
top_words = 5000
(X_train, y_train), (X_test, y_test) = imdb.load_data(nb_words=top_words)
# pad dataset to a maximum review length in words
max_words = 1000
X_train = sequence.pad_sequences(X_train, maxlen=max_words)
X_test = sequence.pad_sequences(X_test, maxlen=max_words)
# create the model
model = Sequential()
model.add(Embedding(top_words, 32, input_length=max_words))
model.add(Convolution1D(nb_filter=32, filter_length=3, border_mode='same',
activation='relu'))
model.add(MaxPooling1D(pool_length=2))
model.add(Flatten())
model.add(Dense(250, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
print(model.summary())
# Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=5, batch_size=128,
verbose=1)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Layer (type)	Output Shape	Param #	Connected to
embedding_7 (Embedding)	(None, 1000, 32)	160000	embedding_input_7[0][0]
convolution1d_4 (Convolution1D)	(None, 1000, 32)	3104	embedding_7[0][0]
maxpooling1d_4 (MaxPooling1D)	(None, 500, 32)	0	convolution1d_4[0][0]
flatten_7 (Flatten)	(None, 16000)	0	maxpooling1d_4[0][0]
dense_13 (Dense)	(None, 250)	4000250	flatten_7[0][0]
dense_14 (Dense)	(None, 1)	251	dense_13[0][0]

Total params: 4163605

None

Train on 25000 samples, validate on 25000 samples

Epoch 1/5

25000/25000 [=====] - 194s - loss: 0.4642 - acc: 0.7551 - val\_loss: 0.4727  
- val\_acc: 0.7914

Epoch 2/5

25000/25000 [=====] - 197s - loss: 0.2475 - acc: 0.9001 - val\_loss: 0.2583  
- val\_acc: 0.8916

Epoch 3/5

25000/25000 [=====] - 186s - loss: 0.1992 - acc: 0.9233 - val\_loss: 0.3620  
- val\_acc: 0.8576

Epoch 4/5

25000/25000 [=====] - 193s - loss: 0.1725 - acc: 0.9344 - val\_loss: 0.2779  
- val\_acc: 0.8884

Epoch 5/5

25000/25000 [=====] - 221s - loss: 0.1475 - acc: 0.9448 - val\_loss: 0.3049  
- val\_acc: 0.8840

Accuracy: 88.40%

In [ ]: